

Spark Operations =

+



TRANSFORMATIONS



ACTIONS



= easy



= medium

Essential Core & Intermediate Spark Operations



General

- map
- filter
- flatMap
- mapPartitions
- mapPartitionsWithIndex
- groupBy
- sortBy

Math / Statistical

- sample
- randomSplit

Set Theory / Relational

- union
- intersection
- subtract
- distinct
- cartesian
- zip

Data Structure / I/O

- keyBy
- zipWithIndex
- zipWithUniqueID
- zipPartitions
- coalesce
- repartition
- repartitionAndSortWithinPartitions
- pipe


-
- reduce
 - collect
 - aggregate
 - fold
 - first
 - take
 - foreach
 - top
 - treeAggregate
 - treeReduce
 - foreachPartition
 - collectAsMap

- count
- takeSample
- max
- min
- sum
- histogram
- mean
- variance
- stdev
- sampleVariance
- countApprox
- countApproxDistinct

- takeOrdered

- saveAsTextFile
- saveAsSequenceFile
- saveAsObjectFile
- saveAsHadoopDataset
- saveAsHadoopFile
- saveAsNewAPIHadoopDataset
- saveAsNewAPIHadoopFile





= easy



= medium

Essential Core & Intermediate PairRDD Operations



General

- flatMapValues
- groupByKey
- reduceByKey
- reduceByKeyLocally
- foldByKey
- aggregateByKey
- sortByKey
- combineByKey

Math / Statistical

- sampleByKey

Set Theory / Relational

- cogroup (=groupWith)
- join
- subtractByKey
- fullOuterJoin
- leftOuterJoin
- rightOuterJoin

Data Structure

- partitionBy

-
- keys
 - values

- countByKey
- countByValue
- countByValueApprox
- countApproxDistinctByKey
- countApproxDistinctByKey
- countByKeyApprox
- sampleByKeyExact






VS




narrow

*each partition of the parent RDD is used by
at most one partition of the child RDD*



wide

*multiple child RDD partitions may depend
on a single parent RDD partition*



LINEAGE

“One of the challenges in providing RDDs as an abstraction is choosing a representation for them that can track lineage across a wide range of transformations.”



“The most interesting question in designing this interface is how to represent dependencies between RDDs.”


“We found it both sufficient and useful to classify dependencies into two types:

- **narrow dependencies**, where each partition of the parent RDD is used by at most one partition of the child RDD
- **wide dependencies**, where multiple child partitions may depend on it.”




narrow


each partition of the parent RDD is used by at most one partition of the child RDD



map, filter




union




join w/ inputs
co-partitioned

wide

multiple child RDD partitions may depend on a single parent RDD partition




groupByKey




join w/ inputs not
co-partitioned





TRANSFORMATIONS




Core Operations



MAP


RDD: ✕





MAP

RDD: `x`




RDD: `y`





MAP

RDD: `x`




RDD: `y`






MAP

RDD: **x**



RDD: **y**





MAP


After `map()` has been applied...

RDD: `x`



before

RDD: `y`




after






MAP



RDD: **y**




Return a new RDD by applying a function to each element of this RDD.






MAP

RDD: **x**



RDD: **y**



`map(f, preservesPartitioning=False)`

Return a new RDD by applying a function to each element of this RDD



```
x = sc.parallelize(["b", "a", "c"])
y = x.map(lambda z: (z, 1))
print(x.collect())
print(y.collect())
```



x: ['b', 'a', 'c']

y: [('b', 1), ('a', 1), ('c', 1)]




```
val x = sc.parallelize(Array("b", "a", "c"))
val y = x.map(z => (z,1))
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```





FILTER


RDD: ✕





FILTER

RDD: `x`




RDD: `y`





FILTER

RDD: x




RDD: y






FILTER

RDD: x



True
emits

RDD: y





After filter() has been applied...


FILTER

RDD: `x`



before

RDD: `y`




after





FILTER



Return a new RDD containing only the elements that satisfy a predicate



```
x = sc.parallelize([1,2,3])
y = x.filter(lambda x: x%2 == 1) #keep odd values
print(x.collect())
print(y.collect())
```



x: [1, 2, 3]

y: [1, 3]




```
val x = sc.parallelize(Array(1,2,3))
val y = x.filter(n => n%2 == 1)
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```





FLATMAP


RDD: `x`






FLATMAP

RDD: `x`




RDD: `y`






FLATMAP

RDD: `x`




RDD: `y`






FLATMAP

RDD: `x`



RDD: `y`





After flatmap() has been applied...


FLATMAP

RDD: `x`



before

RDD: `y`




after






FLATMAP



RDD: `x`



RDD: `y`




Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results






FLATMAP

RDD: **x**



RDD: **y**



`flatMap(f, preservesPartitioning=False)`

Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results



```
x = sc.parallelize([1,2,3])
y = x.flatMap(lambda x: (x, x*100, 42))
print(x.collect())
print(y.collect())
```



```
val x = sc.parallelize(Array(1,2,3))
val y = x.flatMap(n => Array(n, n*100, 42))
println(x.collect().mkString(", "))
println(y.collect().mkString(", "))
```

x: [1, 2, 3]


y: [1, 100, 42, 2, 200, 42, 3, 300, 42]





GROUPBY

RDD: `x`






GROUPBY

RDD: `x`


RDD: `y`






GROUPBY

RDD: **x**




RDD: **y**






GROUPBY

RDD: **x**




RDD: **y**





GROUPBY

RDD: **x**



RDD: **y**





GROUPBY

RDD: **x**



RDD: **y**



`groupBy(f, numPartitions=None)`

Group the data in the original RDD. Create pairs where the key is the output of a user function, and the value is all items for which the function yields this key.



```
x = sc.parallelize(['John', 'Fred', 'Anna', 'James'])
y = x.groupBy(lambda w: w[0])
print [(k, list(v)) for (k, v) in y.collect()]
```



x: ['John', 'Fred', 'Anna', 'James']



```
val x = sc.parallelize(
    Array("John", "Fred", "Anna", "James"))
val y = x.groupBy(w => w.charAt(0))
println(y.collect().mkString(", "))
```


y: [('A', ['Anna']), ('J', ['John', 'James']), ('F', ['Fred'])]





GROUPBYKEY

Pair RDD: `x`






GROUPBYKEY

Pair RDD: `x`


RDD: `y`






GROUPBYKEY

Pair RDD: `x`




RDD: `y`





GROUPBYKEY

RDD: **x**



RDD: **y**



`groupByKey(numPartitions=None)`

Group the values for each key in the original RDD. Create a new pair where the original key corresponds to this collected group of values.

 **x** = sc.parallelize([('B',5),('B',4),('A',3),('A',2),('A',1)])
y = **x**.groupByKey()
print(**x**.collect())
print(list((j[0], list(j[1]))) for j in **y**.collect()))



x: [('B', 5), ('B', 4), ('A', 3), ('A', 2), ('A', 1)]

y: [('A', [2, 3, 1]), ('B', [5, 4])]


 val **x** = sc.parallelize(
 Array(('B',5),('B',4),('A',3),('A',2),('A',1)))
val **y** = **x**.groupByKey()
println(**x**.collect().mkString(", "))
println(**y**.collect().mkString(", "))





UNION



RDD: x



RDD: y

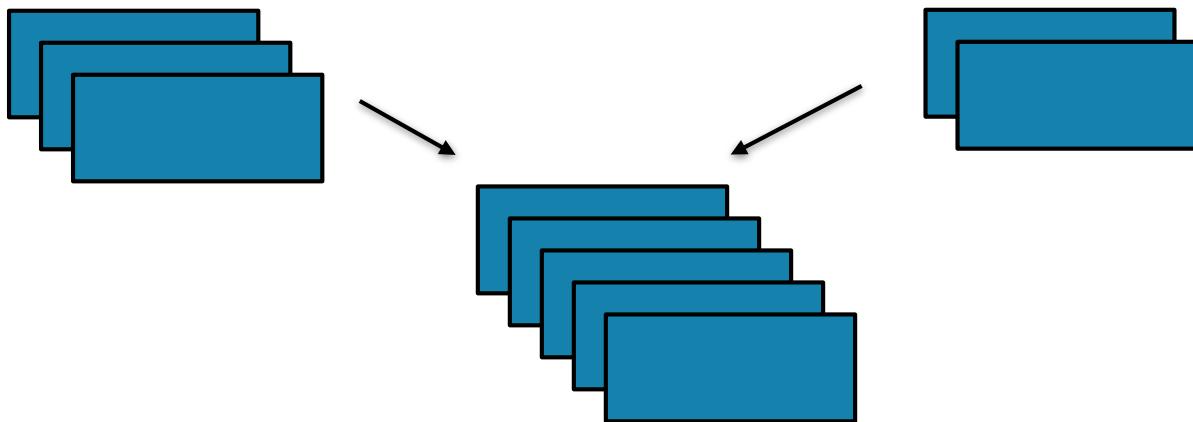


RDD: z





UNION



Return a new RDD containing all items from two original RDDs. Duplicates are *not* culled.

`union(otherRDD)`



```
x = sc.parallelize([1,2,3], 2)
y = sc.parallelize([3,4], 1)
z = x.union(y)
print(z.glom().collect())
```




`x: [1, 2, 3]`

`y: [3, 4]`

`z: [[1], [2, 3], [3, 4]]`



```
val x = sc.parallelize(Array(1,2,3), 2)
val y = sc.parallelize(Array(3,4), 1)
val z = x.union(y)
val zOut = z.glom().collect()
```






JOIN

RDD: `x`




RDD: `y`






JOIN



RDD: x



RDD: y



RDD: z





JOIN



RDD: x



RDD: y



RDD: z







JOIN


RDD: x



RDD: y



JOIN



Return a new RDD containing all pairs of elements having the same key in the original RDDs
`union(otherRDD, numPartitions=None)`




```
x = sc.parallelize([("a", 1), ("b", 2)])
y = sc.parallelize([("a", 3), ("a", 4), ("b", 5)])
z = x.join(y)
print(z.collect())
```



```
x: [("a", 1), ("b", 2)]
y: [("a", 3), ("a", 4), ("b", 5)]
z: [('a', (1, 3)), ('a', (1, 4)), ('b', (2, 5))]
```





```
val x = sc.parallelize(Array(("a", 1), ("b", 2)))
val y = sc.parallelize(Array(("a", 3), ("a", 4), ("b", 5)))
val z = x.join(y)
println(z.collect().mkString(", "))
```





DISTINCT


RDD: `x`







DISTINCT

RDD: \mathbf{x}




RDD: \mathbf{y}







DISTINCT

RDD: `x`




RDD: `y`





DISTINCT



Return a new RDD containing distinct items from the original RDD (omitting all duplicates)

`distinct(numPartitions=None)`



```
x = sc.parallelize([1,2,3,3,4])
y = x.distinct()

print(y.collect())
```




`x: [1, 2, 3, 3, 4]`

`y: [1, 2, 3, 4]`



```
val x = sc.parallelize(Array(1,2,3,3,4))
val y = x.distinct()


println(y.collect().mkString(", "))
```





COALESCE


RDD: **x**






COALESCE

RDD: **x**




RDD: **y**






COALESCE

RDD: **x**




RDD: **y**





COALESCE



Return a new RDD which is reduced to a smaller number of partitions

`coalesce(numPartitions, shuffle=False)`



```
x = sc.parallelize([1, 2, 3, 4, 5], 3)
y = x.coalesce(2)
print(x.glom().collect())
print(y.glom().collect())
```



`x: [[1], [2, 3], [4, 5]]`

`y: [[1], [2, 3, 4, 5]]`




```
val x = sc.parallelize(Array(1, 2, 3, 4, 5), 3)
val y = x.coalesce(2)
val xOut = x.glom().collect()
val yOut = y.glom().collect()
```





KEYBY

RDD: **x**




RDD: **y**





KEYBY

RDD: **x**




RDD: **y**






KEYBY

RDD: **x**




RDD: **y**






KEYBY

RDD: **x**




RDD: **y**






KEYBY

RDD: **x**



RDD: **y**



keyBy(*f*)

Create a Pair RDD, forming one pair for each item in the original RDD. The pair's key is calculated from the value via a user-supplied function.



```
x = sc.parallelize(['John', 'Fred', 'Anna', 'James'])
y = x.keyBy(lambda w: w[0])
print y.collect()
```



x: ['John', 'Fred', 'Anna', 'James']

y: [('J','John'), ('F','Fred'), ('A','Anna'), ('J','James')]




```
val x = sc.parallelize(
    Array("John", "Fred", "Anna", "James"))
val y = x.keyBy(w => w.charAt(0))
println(y.collect().mkString(", "))
```







ZIP

RDD: x




RDD: y






ZIP



RDD: x



RDD: y




RDD: z






ZIP



RDD: x



RDD: y



RDD: z






ZIP



RDD: x



RDD: y




RDD: z





ZIP



Return a new RDD containing pairs whose key is the item in the original RDD, and whose value is that item's corresponding element (same partition, same index) in a second RDD

`zip(otherRDD)`



```
x = sc.parallelize([1, 2, 3])
y = x.map(lambda n:n*n)
z = x.zip(y)

print(z.collect())
```



`x: [1, 2, 3]`


`y: [1, 4, 9]`


`z: [(1, 1), (2, 4), (3, 9)]`




```
val x = sc.parallelize(Array(1,2,3))
val y = x.map(n=>n*n)
val z = x.zip(y)

println(z.collect().mkString(", "))
```





ACTIONS



Core Operations



VS

distributed


occurs across the cluster

driver


result must fit in driver JVM



GETNUMPARTITIONS



GETNUMPARTITIONS



2

getNumPartitions()

Return the number of partitions in RDD



```
x = sc.parallelize([1,2,3], 2)
y = x.getNumPartitions()

print(x.glom().collect())
print(y)
```



x: [[1], [2, 3]]



y: 2





```
val x = sc.parallelize(Array(1,2,3), 2)
val y = x.partitions.size
val xOut = x.glom().collect()
println(y)
```



COLLECT



COLLECT



`collect()`

Return all items in the RDD to the driver in a single list



```
x = sc.parallelize([1,2,3], 2)
y = x.collect()

print(x.glom().collect())
print(y)
```



`x: [[1], [2, 3]]`

`y: [1, 2, 3]`




```
val x = sc.parallelize(Array(1,2,3), 2)
val y = x.collect()


val xOut = x.glom().collect()
println(y)
```




REDUCE




REDUCE



REDUCE



REDUCE



`reduce(f)`

Aggregate all the elements of the RDD by applying a user function pairwise to elements and partial results, and returns a result to the driver



```
x = sc.parallelize([1,2,3,4])
y = x.reduce(lambda a,b: a+b)

print(x.collect())
print(y)
```



`x:` [1, 2, 3, 4]
`y:` 10





```
val x = sc.parallelize(Array(1,2,3,4))
val y = x.reduce((a,b) => a+b)


println(x.collect.mkString(", "))
println(y)
```



MAX



MAX



`max()`

Return the maximum item in the RDD



```
x = sc.parallelize([2,4,1])
y = x.max()

print(x.collect())
print(y)
```




`x: [2, 4, 1]`

`y: 4`





```
val x = sc.parallelize(Array(2,4,1))
val y = x.max



println(x.collect().mkString(", "))
println(y)
```



SUM



SUM



sum()

Return the sum of the items in the RDD



```
x = sc.parallelize([2,4,1])
y = x.sum()

print(x.collect())
print(y)
```




x: [2, 4, 1]

y: 7




```
val x = sc.parallelize(Array(2,4,1))
val y = x.sum


println(x.collect().mkString(", "))
println(y)
```




MEAN



2.33333333



MEAN



2.333333


`mean()`

Return the mean of the items in the RDD



```
x = sc.parallelize([2,4,1])
y = x.mean()

print(x.collect())
print(y)
```




`x:` [2, 4, 1]

`y:` 2.333333




```
val x = sc.parallelize(Array(2,4,1))
val y = x.mean


println(x.collect().mkString(", "))
println(y)
```




STDEV



1.2472191



STDEV



1.2472191



A screenshot of a terminal window displaying the numerical value 1.2472191, which is the standard deviation of the RDD.

stdev()

Return the standard deviation of the items in the RDD



```
x = sc.parallelize([2,4,1])
y = x.stdev()

print(x.collect())
print(y)
```




x: [2, 4, 1]

y: 1.2472191



```
val x = sc.parallelize(Array(2,4,1))
val y = x.stdev


println(x.collect().mkString(", "))
println(y)
```




COUNTBYKEY



{'A': 1, 'J': 2, 'F': 1}



COUNTBYKEY



countByKey()

Return a map of keys and counts of their occurrences in the RDD



```
x = sc.parallelize([('J', 'James'), ('F', 'Fred'),  
                    ('A', 'Anna'), ('J', 'John')])
```

```
y = x.countByKey()  
print(y)
```




```
val x = sc.parallelize(Array(('J', "James"), ('F', "Fred"),  
                           ('A', "Anna"), ('J', "John")))
```


```
val y = x.countByKey()  
println(y)
```

x: [('J', 'James'), ('F', 'Fred'),
 ('A', 'Anna'), ('J', 'John')]


y: {'A': 1, 'J': 2, 'F': 1}



SAVE AS TEXT FILE



SAVEASTEXTFILE



`saveAsTextFile(path, compressionCodecClass=None)`

Save the RDD to the filesystem indicated in the path



```
dbutils.fs.rm("/temp/demo", True)
x = sc.parallelize([2,4,1])
x.saveAsTextFile("/temp/demo")
```

```
y = sc.textFile("/temp/demo")
print(y.collect())
```



`x: [2, 4, 1]`

`y: [u'2', u'4', u'1']`



```
dbutils.fs.rm("/temp/demo", true)
val x = sc.parallelize(Array(2,4,1))
x.saveAsTextFile("/temp/demo")
```

```
val y = sc.textFile("/temp/demo")
println(y.collect().mkString(", "))
```



LAB



Q&A

